

Specifying the roundoff errors of SONNX components

Franck Vedrine – CEA
with the help of Eric Jenn – IRT Saint-Exupéry



Accumulation of roundoff errors for SONNX tensor algorithms

- Numerous possible numerical format for computations
 - floating-point: double, float, bfloat16, ...
 - fixed-point: different sizes for the fractional part
- Numerous possible algorithms
 - Ex: optimizations for SoftMax
 - With Error Free Transformation (Rump & all) techniques to gain precision
 - Ex: an EFT for Addition: FastTwoSum, $a + b = s + t$

```
void FastTwoSum(double a, double b, double* s, double* t)
{ // No unsafe optimizations!
    *s = a + b ;
    *t = b - (* s - a ) ;
}
```

- How to define **precise** and **concise** accuracy formula that are **verifiable for each implementation?**



Available technologies

- Arithmetic for the accumulation of roundoff error
 - Concrete comparison/difference between the implementation format and a better data format
 - Stochastic arithmetic – Cadna (source instrumentation - LIP6), Verrou (binary instrumentation with valgrind - EDF), Verificarlo (intermediary instrumentation with llvm - UVSQ)
 - Affine arithmetic – Fluctuat (abstract interpreter - CEA), FLDLib (source instrumentation, “abstract compiler” - CEA) – soundness
 - Formal arithmetic and simplification – soundness
- in different analysis technologies
 - Instrumentation of floating-point data
 - Abstract interpretation – all numerical data
 - Instrumentation of all numerical data



Inference rules of atomic operations

- Every concrete value f is defined as $x + e$ where
 - x is a real number that would be computed if the numerical format was ideal
 - e is a real number representing the numerical error that is the difference between f the floating-point or fixed-point value and x
 - $+, -, \times, /$ are operations in real numbers, $+_{\text{impl}}, -_{\text{impl}}, *_{\text{impl}}/_{\text{impl}}$ are operations in floating/fixed-point
- The **induced error** is a function taking an input error e and returning how it is amplified
- The **introduced error** is a new error that the operation creates to approximate the ideal result in a storing objective

With these definitions, the inference rules are

- $-_{\text{impl}}(x + e) = -x + -e$ is an exact operation
- $(x_0 + e_0) +_{\text{impl}} (x_1 + e_1) = (x_0 + x_1) + (e_0 + e_1) + \text{ introduced error of } ((x_0 + x_1) + (e_0 + e_1))$
- $(x_0 + e_0) -_{\text{impl}} (x_1 + e_1) = (x_0 - x_1) + (e_0 - e_1) + \text{ introduced error of } ((x_0 - x_1) + (e_0 - e_1))$
- $(x_0 + e_0) *_{\text{impl}} (x_1 + e_1) = (x_0 \times x_1) + (e_0 \times x_1 + e_1 \times (x_0 + e_0)) + \text{ introduced error of } ((x_0 \times x_1) + (e_0 \times x_1 + e_1 \times (x_0 + e_0)))$
- $(x_0 + e_0) /_{\text{impl}} (x_1 + e_1) = \frac{x_0}{x_1} + \frac{x_0 \times e_1 - e_0 \times x_1}{x_1 \times (x_1 + e_1)} + \text{ introduced error of } \left(\frac{x_0}{x_1} + \frac{x_0 \times e_1 - e_0 \times x_1}{x_1 \times (x_1 + e_1)}\right)$



Focus on the introduced error

- In floating-point, the introduced error of (x)
 - $\leq |x| \times u$, where u is the semi-ulp of 1, that is 2^{-53} for the double type
 - could be improved with the Sterbenz Lemma $|a| -_{\text{impl}} |b|$ is exact if $\frac{1}{2}|a| \leq |b| \leq 2|a|$
 - this lemma needs the operation and a range for the operands
 - is deterministic in every mode and can be exactly computed if x is known (see EFT)
- In fixed-point, the introduced error of (x)
 - It often needs the operations and the operands
 - $x +_{\text{impl}} y$ does not introduce any error
 - $\leq u$, where u is given by the fractional part of the fixed-point format



Objectives for SONNX components

1. Write an accuracy specification for tensor operations that are verifiable
 - Tensor operations combine an imprecise (potentially unbound) number of atomic operations (dimensions)
 - The specification should allow multiple implementations
2. Check if an implementation verifies a specification with analyses ensuring some guarantees
3. Show that an implementation is better than another one from the accuracy point of view



Subjective approach/methodology

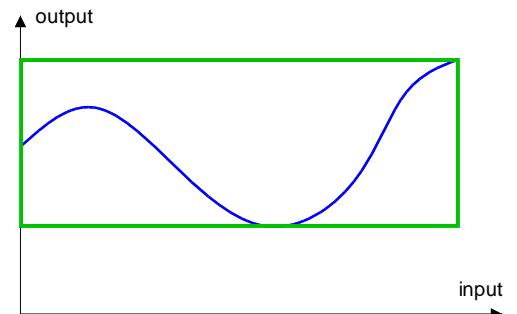
- 1. Write an accuracy specification for tensor operations that are verifiable**
 1. Write 2 or 3 (more or less approximated) specifications with pen and paper for low dimensions
 - then extrapolate specifications w.r.t. a formal variable representing the tensor dimension
 - and finally prove by induction that the specification is consistent for the component
 2. or generate the specification with symbolic formula of slide 4
 - using some “widening” operation (Abstract Interpretation wording) over the tensor dimension
- 2. Check if an implementation verifies a specification with analyses ensuring some guarantees**
 - Check with stochastic arithmetic on some representative use-cases
 - Check with affine arithmetic and generate accuracy bounds on some representative numerical scenarios = use-cases with input ranges
 - Be ready to subdivide non-linear components
 - Check with symbolic arithmetic and generate accuracy formula
 - Be ready to guide the simplifications of the symbolic abstract terms
- 3. Show that an implementation is better than another one from the accuracy point of view**
 - Compare the accuracy bounds/formula generated by the analyses

Symbolic arithmetic and affine arithmetic

- Symbolic arithmetic

- $(x_0 + e_0) + \text{impl } (x_1 + e_1) = (x_0 + x_1) + (\mathbf{e_0} + \mathbf{e_1}) + ((x_0 + x_1) + (e_0 + e_1)) \times [-u, u]$
- $(x_0 + e_0) * \text{impl } (x_1 + e_1) = (x_0 \times x_1) + (\mathbf{e_0} \times \mathbf{x_1} + \mathbf{e_1} \times (\mathbf{x_0} + \mathbf{e_0}))$
 $+ ((x_0 \times x_1) + (e_0 \times x_1 + e_1 \times (x_0 + e_0))) \times [-u, u]$
- $(x_0 + e_0) / \text{impl } (x_1 + e_1) = \frac{x_0}{x_1} + \frac{\mathbf{x_0} \times \mathbf{e_1} - \mathbf{e_0} \times \mathbf{x_1}}{\mathbf{x_1} \times (\mathbf{x_1} + \mathbf{e_1})} + \left(\frac{x_0}{x_1} + \frac{x_0 \times e_1 - e_0 \times x_1}{x_1 \times (x_1 + e_1)} \right) \times [-u, u]$

- The combination of operation \Rightarrow big terms
- No quantitative results
- Exact formula, no over-approximation
- Need some interval abstractions and simplifications to find generic formula



- Affine arithmetic

- Normalization for non-linear operations \Rightarrow some over-approximations (often too imprecise for large input intervals)
- Quantitative results and intervals for the error
- Limitation of the number of shared symbols to find generic formula

Specificities of the arithmetic with affine forms

- The ideal value and the error (x, e) are represented by $\alpha_0 + \alpha_1 \times \varepsilon_1 + \alpha_2 \times \varepsilon_2 + \alpha_3 \times \varepsilon_3 + \dots$ where
 - $(\alpha_i)_{i \in [0,n]}$ are concrete values $\in \mathbb{R}^{n+1}$, $(\varepsilon_i)_{i \in [1,n]}$ are symbolic shared variables $\in [-1, 1]^n$
 - Hence, $x \in [1.0, 6.0]$ becomes $x = 3.5 + 2.5 \times \varepsilon_1$, with $\varepsilon_1 = \frac{x-3.5}{2.5} \in [-1, 1]$
- $(x_0 + e_0) * \text{impl} (x_1 + e_1) = (x_0 \times x_1) + (\mathbf{e}_0 \times \mathbf{x}_1 + \mathbf{e}_1 \times (x_0 + e_0))$

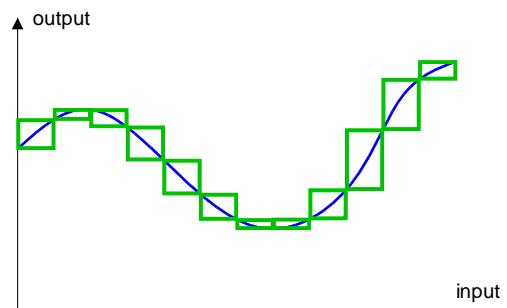
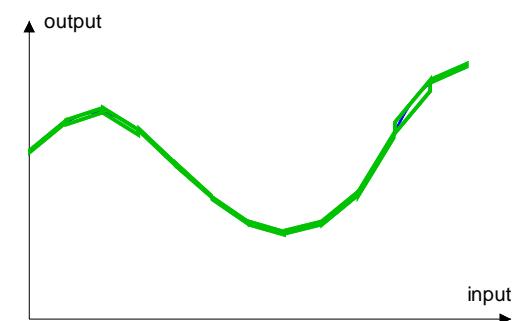
+ introduced error of $((x_0 \times x_1) + (\mathbf{e}_0 \times \mathbf{x}_1 + \mathbf{e}_1 \times (x_0 + e_0)))$

$$(\alpha_0 + \alpha_1 \times \varepsilon_1 + \alpha_2 \times \varepsilon_2) + (\beta_0 + \beta_1 \times \varepsilon_1 + \beta_2 \times \varepsilon_2) = (\alpha_0 + \beta_0) + (\alpha_1 + \beta_1) \times \varepsilon_1 + (\alpha_2 + \beta_2) \times \varepsilon_2, \text{ no additional over-approximation}$$

$$(\alpha_0 + \alpha_1 \times \varepsilon_1 + \alpha_2 \times \varepsilon_2) \times (\beta_0 + \beta_1 \times \varepsilon_1 + \beta_2 \times \varepsilon_2) = ((\alpha_0 \times \beta_0) + \frac{\alpha_1 \times \beta_1}{2} + \frac{\alpha_2 \times \beta_2}{2}) + (\alpha_1 \times \beta_0 + \beta_1 \times \alpha_0) \times \varepsilon_1 + (\alpha_2 \times \beta_0 + \beta_2 \times \alpha_0) \times \varepsilon_2 + \frac{\alpha_1 \times \beta_1}{2} \times \mu_1 + \frac{\alpha_2 \times \beta_2}{2} \times \mu_2 + (\alpha_1 \times \beta_2 + \alpha_2 \times \beta_1) \mu_{12}$$

- over-approximation: $\mu_i = 2 \times \varepsilon_i^2 - 1$ is a new fresh symbolic variable $\in [-1, 1]$
- over-approximation: $\mu_{ij} = \varepsilon_i \times \varepsilon_j$ is a new fresh symbolic variable $\in [-1, 1]$

- The simplification under the “normal” affine form is automatic
- The range of $\alpha_0 + \sum_{i=1}^n \alpha_i \times \varepsilon_i$ is $[\alpha_0 - \sum_{i=1}^n |\alpha_i|, \alpha_0 + \sum_{i=1}^n |\alpha_i|]$
- Non-linear operations generate over-approximations
⇒ subdividing the input ranges improves the quality of the results





Example of Matrix multiplication

- Implicit need of a model of matrix multiplication with a numerical error containing all possible implementations
- Symbolic abstract inference of the error for this model with the rules of slide 4

Let us denote $|X| = \max_{i,j} |X_{ij}|$ and $|E_X| = \max_{i,j} |E_{xij}|$

- case $\dim_X = (1,1)$, $\dim_Y = (1,1)$, $(X) = (x_0 + e_{x0})$, $(Y) = (y_0 + e_{y0})$
 - Symbolic induced error: $e_{y0} \times x_0 + e_{x0} \times (y_0 + e_{y0})$
 - Approximate induced error:
 - $[-|X|, |X|] \times e_{y0} + [-|Y|, |Y|] \times e_{x0} + e_{x0} \times e_{y0}$
 - $[-|X| \times |E_Y| - (|Y| + |E_Y|) \times |E_X|, |X| \times |E_Y| + (|Y| + |E_Y|) \times |E_X|]$
 - Introduced error (*double*, and $|X|, |Y|$ normal floating-point):
 - $[-u, u] \times ((x_0 + e_{x0}) \times (y_0 + e_{y0}))$
 - $[-u, u] \times (|X| + |E_X|) \times (|Y| + |E_Y|)$



Example of Matrix multiplication (2)

Let us denote $|X| = \max_{i,j} |X_{ij}|$ and $|E_X| = \max_{i,j} |E_{xij}|$

- case $\dim_X = (1,2)$, $\dim_Y = (2,1)$, $(X) = (x_0 + e_{x0}, x_1 + e_{x1})$, $(Y) = \begin{pmatrix} y_0 + e_{y0} \\ y_1 + e_{y1} \end{pmatrix}$
 - Symbolic induced error: $e_{y0} \times x_0 + e_{x0} \times (y_0 + e_{y0}) + e_{y1} \times x_1 + e_{x1} \times (y_1 + e_{y1})$
 - Approximate induced error:
 - $[-|X|, |X|] \times (e_{y0} + e_{y1}) + [-|Y|, |Y|] \times (e_{x0} + e_{x1}) + e_{x0} \times e_{y0} + e_{x1} \times e_{y1}$
 - $2 \times [-|X| \times |E_Y| - (|Y| + |E_Y|) \times |E_X|, |X| \times |E_Y| + (|Y| + |E_Y|) \times |E_X|]$
 - Introduced error:
 - $$\begin{aligned} & \left((1 + [-u, u]) \times \left((x_0 + e_{x0}) \times (y_0 + e_{y0}) \right) + (1 + [-u, u]) \times \left((x_1 + e_{x1}) \times (y_1 + e_{y1}) \right) \right) \times (1 + [-u, u]) \\ & \quad - \left((x_0 + e_{x0}) \times (y_0 + e_{y0}) + (x_1 + e_{x1}) \times (y_1 + e_{y1}) \right) \\ & = \left((x_0 + e_{x0}) \times (y_0 + e_{y0}) + (x_1 + e_{x1}) \times (y_1 + e_{y1}) \right) \times ((1 + [-u, u])^2 - 1) \end{aligned}$$
 - $(2 + [-u, u]) \times [-u, u] \times (|X| + |E_X|) \times (|Y| + |E_Y|)$



Example of Matrix multiplication (3)

Let us denote $|X| = \max_{i,j} |X_{ij}|$ and $|E_X| = \max_{i,j} |E_{xij}|$

- case $\dim_X = (1,3), \dim_Y = (3,1), (X) = (x_0 + e_{x0}, x_1 + e_{x1}, x_2 + e_{x2}), (Y) = \begin{pmatrix} y_0 + e_{y0} \\ y_1 + e_{y1} \\ y_2 + e_{y2} \end{pmatrix}$
- Symbolic induced error: $\sum_{i=0}^2 (e_{yi} \times x_i + e_{xi} \times (y_i + e_{yi}))$
- Approximate induced error:
 - $\sum_{i=0}^2 (e_{yi} \times [-|X|, |X|]) + \sum_{i=0}^2 (e_{xi} \times [-|Y|, |Y|]) + \sum_{i=0}^2 (e_{xi} \times e_{yi})$
 - $3 \times [-|X| \times |E_Y| - (|Y| + |E_Y|) \times |E_X|, |X| \times |E_Y| + (|Y| + |E_Y|) \times |E_X|]$
- Introduced error:
 - $\left(\sum_{i=0}^2 ((x_i + e_{xi}) \times (y_i + e_{yi})) \times ((1 + [-u, u])^{4-i} - 1) \right)$ implementation associativity dependency
 - $\left[n - (1+u)^2 \times \frac{(1+u)^3 - 1}{u}, (1+u)^2 \times \frac{(1+u)^3 - 1}{u} - n \right] \times (|X| + |E_X|) \times (|Y| + |E_Y|)$
less precise formula, but independent from any associativity rule



Example of Matrix multiplication (3) - extrapolation

- case $\dim_X = (1, n), \dim_Y = (n, 1)$
 - Symbolic induced error: $\sum_{i=0}^{n-1} (e_{yi} \times x_i + e_{xi} \times (y_i + e_{yi}))$
 - Approximate induced error:
 - $\sum_{i=0}^{n-1} (e_{yi} \times [-|X|, |X|]) + \sum_{i=0}^{n-1} (e_{xi} \times [-|Y|, |Y|]) + \sum_{i=0}^{n-1} (e_{xi} \times e_{yi})$
 - $n \times [-|X| \times |E_Y| - (|Y| + |E_Y|) \times |E_X|, |X| \times |E_Y| + (|Y| + |E_Y|) \times |E_X|]$
 - Introduced error:
 - $(\sum_{i=0}^{n-1} ((x_i + e_{xi}) \times (y_i + e_{yi}))) \times ((1 + [-u, u])^{n+1-i} - 1)$
 - $[n - (1 + u)^2 \times \frac{(1+u)^{n-1}}{u}, (1 + u)^2 \times \frac{(1+u)^{n-1}}{u} - n] \times (|X| + |E_X|) \times (|Y| + |E_Y|)$
- Provable by induction over n – if your model of matrix multiplication is inductively defined over n

Example of Matrix multiplication (3) - generalization

- case $\dim_X = (p, n), \dim_Y = (n, q)$
 - Symbolic induced error for z_{ij} : $\sum_{k=0}^{n-1} (e_{ykj} \times x_{ik} + e_{xik} \times (y_{kj} + e_{ykj}))$
 - Approximate induced error:
 - $n \times [-|X| \times |E_Y| - (|Y| + |E_Y|) \times |E_X|, |X| \times |E_Y| + (|Y| + |E_Y|) \times |E_X|]$
 - $\sum_{k=0}^{n-1} (e_{ykj} \times [-|X|, |X|]) + \sum_{k=0}^{n-1} (e_{xik} \times [-|Y|, |Y|]) + \sum_{k=0}^{n-1} (e_{xik} \times e_{ykj})$
 - Introduced error:
 - $\left(\sum_{k=0}^{n-1} ((x_{ik} + e_{xik}) \times (y_{kj} + e_{ykj})) \right) \times \left((1 + [-u, u])^{n+1-i} - 1 \right)$
 - $\left[n - (1+u)^2 \times \frac{(1+u)^n - 1}{u}, (1+u)^2 \times \frac{(1+u)^n - 1}{u} - n \right] \times (|X| + |E_X|) \times (|Y| + |E_Y|)$
- Provable by induction over p, q



Checking an implementation

- Ideally, use a symbolic abstract static analysis that has
 - simplifications over symbolic terms
 - a widening operator able to extrapolate formula and to prove them by induction
 - If it uses double ($u' = 2^{-53}$) instead of float ($u = 2^{-24}$), it should prove
 - $\forall e_{ykJ}, e_{xik}. \left| \sum_{k=0}^{n-1} (e_{ykJ} \times [-|X|, |X|]) + \sum_{k=0}^{n-1} (e_{xik} \times [-|Y|, |Y|]) + \sum_{k=0}^{n-1} (e_{xik} \times e_{ykJ}) \right| \leq n \times [|X| \times |E_Y| + (|Y| + |E_Y|) \times |E_X|]$
 - $\forall e_{ykJ}, e_{xik}. \left| \sum_{k=0}^{n-1} ((x_{ik} + e_{xik}) \times (y_{kj} + e_{ykJ})) \times ((1 + [-u', u'])^{n+1-i} - 1) \right| \leq \left((1+u)^2 \times \frac{(1+u)^n - 1}{u} - n \right) \times (|X| + |E_X|) \times (|Y| + |E_Y|)$
- or $(1+u')^2 \times \frac{(1+u')^n - 1}{u'} \leq (1+u)^2 \times \frac{(1+u)^n - 1}{u}$ for any n

Checking an implementation – more automatic

- Define a numerical scenario:
 - all coefficients are in $[-1, 1]$,
 - choose $n = 6$
 - the absolute numerical error of the input coefficients is bound by 10^{-6}
- The static analysis (with FLDLib for instance) should be able to prove in `double` that the final error e
 - $|e| \leq 1.2 \times 10^{-5}$ (induced error) + $\left((1 + 2^{-53})^2 \times \frac{(1+2^{-53})^6 - 1}{2^{-53}} - 6 \right) \times (1 + 10^{-6})^2$ (introduced error)
 - no subdivision is required since the worst case has maximal coefficients for both matrices



From the implementation point of view

- When it computes $\sum_{k=0}^{n-1} x_{ik} \times y_{kj}$, it may compute $\left(\sum_{k=0}^{(n-1)/2} x_{ik} \times y_{kj}\right) + \left(\sum_{k=(n+1)/2}^{n-1} x_{ik} \times y_{kj}\right)$
- Hence for $n = 4$, it can claim that the introduced error is bound by

$$4 \times ((1+u)^3 - 1) \sim 7.153 \times 10^{-7} \text{ for } u = 2^{-24} \text{ the float semi-ulp}$$

which is better than the specification

$$(1+u)^2 \times \frac{(1+u)^4 - 1}{u} - 4 \sim 8.35 \times 10^{-7} \text{ for } u = 2^{-24} \text{ the float semi-ulp}$$

- It can also compare to another implementation



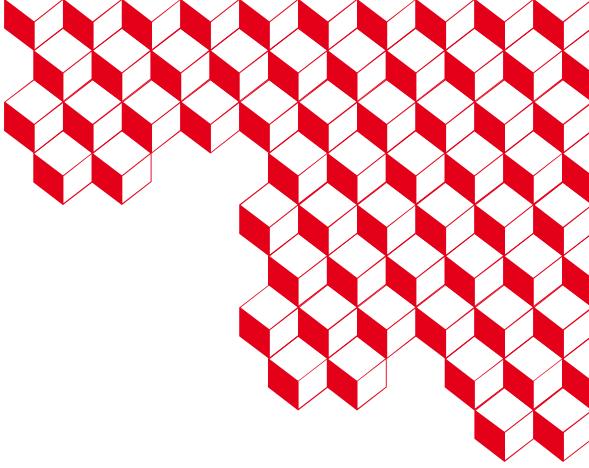
Work in Progress

- Writing models of SONNX components to be able to perform proof by induction
- Writing formula for SONNX components – ex convolution (model conv2d.c)
 - for the induced numerical error
 - for the introduced numerical error
- Writing symbolic abstract domains to synthesize accuracy formula with widening operator implementing proof by induction
 - Simplification of symbolic expressions
 - Extrapolation
 - Proof by induction
- Any help is welcome !



Conclusion – no silver bullet

- The accuracy formula will be approximated, so several specifications will be proposed depending on the approximation level
- The accuracy formula will be complex
 - but it is possible to show how they are obtained with the different steps of the methodology
 - these steps use induction over tensor dimension, so the more complex formula may be inductively defined w.r.t this dimension – we need a program to check them
- An implementation can check the specification
 - with test-cases in stochastic arithmetic
 - precise, fully automatic, realistic – standard deviation, but low guarantees
 - with numerical scenario in affine arithmetic
 - some formal guarantee, less precise and it may need annotations to subdivide
 - in symbolic arithmetic
 - strong guarantee, but induction proofs to fully automatize and certificate to produce (in why3?)
 - If the implementation use EFT to gain accuracy, it should explain it to the analysis



Thanks!